

Gambling Prisoner

Sam Reeves

Smith is in jail and has 1 dollar; he can get out on bail if he has 8 dollars. A guard agrees to make a series of bets with him. If Smith bets A dollars, he wins A dollars with probability .4 and loses A dollars with probability .6.

Find the probability that he wins 8 dollars before losing all of his money if:

(a) he bets 1 dollar each time (timid strategy).

We can analyze this as an absorbing markov chain.

```
library(markovchain)
p <- 0.4
q <- 1 - p
```

For a transition matrix P in canonical form:

$$P = \begin{pmatrix} I & R \\ 0 & Q \end{pmatrix}$$
$$N = (I - Q)^{-1}$$

```
fundamental_matrix <- function(markov_chain_object) {
  # N = (I - Q)^-1
  m <- markov_chain_object

  r <- length(unlist(absorbingStates(m)))
  t <- length(unlist(transientStates(m)))

  P <- as(canonicForm(m), "matrix")
  I <- diag(t)

  Q <- P[(r+1):(r+t), (r+1):(r+t)]
  N <- solve(I - Q)
  return(list(N, P, r))
}
```

$$B = NR$$

```
absorbing_chain_probabilities <- function(m) {
  # B = NR
  x <- fundamental_matrix(m)
  N <- x[[1]]
  P <- x[[2]]

  r <- as.numeric(x[[3]])
  t <- dim(N)[1]

  R <- P[(r+1):(r+t), 1:r]
  B <- N %*% R

  return(B)
}
```

```
m_a <- matrix(c(1, rep(0, 8),
               q, 0, p, rep(0, 6),
               0, q, 0, p, rep(0, 5),
               rep(0, 2), q, 0, p, rep(0, 4),
               rep(0, 3), q, 0, p, rep(0, 3),
               rep(0, 4), q, 0, p, rep(0, 2),
               rep(0, 5), q, 0, p, 0,
               rep(0,6), q, 0, p,
               rep(0,8), 1), nrow = 9, byrow = TRUE,
             dimnames = list(c(0:8), c(0:8)))
```

```
(mc_a <- new('markovchain', transitionMatrix = m_a, states = colnames(m_a)))
```

```
## Unnamed Markov chain
## A 9 - dimensional discrete Markov Chain defined by the following states:
## 0, 1, 2, 3, 4, 5, 6, 7, 8
## The transition matrix (by rows) is defined as follows:
##   0  1  2  3  4  5  6  7  8
## 0 1.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
## 1 0.6 0.0 0.4 0.0 0.0 0.0 0.0 0.0
## 2 0.0 0.6 0.0 0.4 0.0 0.0 0.0 0.0
## 3 0.0 0.0 0.6 0.0 0.4 0.0 0.0 0.0
## 4 0.0 0.0 0.0 0.6 0.0 0.4 0.0 0.0
## 5 0.0 0.0 0.0 0.0 0.6 0.0 0.4 0.0
## 6 0.0 0.0 0.0 0.0 0.0 0.6 0.0 0.4
## 7 0.0 0.0 0.0 0.0 0.0 0.0 0.6 0.4
## 8 0.0 0.0 0.0 0.0 0.0 0.0 0.0 1.0
```

```
#absorbing_chain_probabilities(mc_a)
#plot(absorbing_chain_probabilities(mc_a)[,2],
#      type='S', main='Success Rate', xlab='', ylab='')
```

	0	8
1	0.9796987	0.02030135
2	0.9492466	0.05075337
3	0.9035686	0.09643140
4	0.8350515	0.16494845
5	0.7322760	0.26772403
6	0.5781126	0.42188739
7	0.3468676	0.65313243




Figure 1: Timid betting strategy

The probability of getting to \$8 from \$1 with this strategy is ≈ 0.0203 .

(b) he bets, each time, as much as possible but not more than necessary to bring his fortune up to 8 dollars (bold strategy).

```
m_b <- matrix(c(1, rep(0,8),
               q, 0, p, rep(0,6),
               q, rep(0,3), p, rep(0,4),
               q, rep(0,5), p, rep(0,2),
               q, rep(0,7), p,
               rep(0,2), q, rep(0,5), p,
               rep(0,4), q, rep(0,3), p,
               rep(0,6), q, 0, p,
               rep(0,8), 1),
             nrow = 9, byrow = TRUE,
             dimnames = list(c(0:8), c(0:8)))

(mc_b <- new('markovchain', transitionMatrix = m_b, states = colnames(m_b)))
```

```
## Unnamed Markov chain
## A 9 - dimensional discrete Markov Chain defined by the following states:
## 0, 1, 2, 3, 4, 5, 6, 7, 8
## The transition matrix (by rows) is defined as follows:
##   0 1  2 3  4 5  6 7  8
## 0 1.0 0 0.0 0 0.0 0 0.0 0 0.0
## 1 0.6 0 0.4 0 0.0 0 0.0 0 0.0
## 2 0.6 0 0.0 0 0.4 0 0.0 0 0.0
## 3 0.6 0 0.0 0 0.0 0 0.4 0 0.0
## 4 0.6 0 0.0 0 0.0 0 0.0 0 0.4
## 5 0.0 0 0.6 0 0.0 0 0.0 0 0.4
## 6 0.0 0 0.0 0 0.6 0 0.0 0 0.4
## 7 0.0 0 0.0 0 0.0 0 0.6 0 0.4
## 8 0.0 0 0.0 0 0.0 0 0.0 0 1.0
```

```
#absorbing_chain_probabilities(mc_b)
#plot(absorbing_chain_probabilities(mc_b)[,2], type='S', main='Success Rate', xlab='', ylab='')
```

	0	8
1	0.936	0.064
2	0.840	0.160
3	0.744	0.256
4	0.600	0.400
5	0.504	0.496
6	0.360	0.640
7	0.216	0.784




Figure 2: Bold betting strategy

The probability of getting to \$8 from \$1 with this strategy is ≈ 0.064 .

(c) Which strategy gives Smith the better chance of getting out of jail?

The adaptive strategy gives a better chance from every starting state.

NOTE: `solve(I-Q)` causes the knit function to hang... However, you can run the code from the .Rmd file after uncommenting the lines above the figures.